ASSIGNMENT - 4

NAME :DEVIKA.M

REG NO : 611419205006

SUB : ARTIFICIAL INTELLIGENCE

COLLEGE : MAHENDRA ENGINEERING COLLEGE FOR WOMEN

SMS Spam Prediction using Naive Bayes

Hello!

In this notebook I am building a Naive Bayes classifier, to identify SMS spam from non-spam.

I am going to compare the Bernoulli and the Multinomial NB classifiers.

Table of contents

Libraries and dataset

This Python 3 environment comes with many helpful analytics libraries installed

```
# It is defined by the kaggle/python Docker image: https://github.com/kaggle/docker-python
# For example, here's several helpful packages to load


import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt # data visualization purposes
import seaborn as sns # statistical data visualization
%matplotlib inline


# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list all files under the input directory


import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))


# You can write up to 20GB to the current directory (/kaggle/working/) that gets preserved as output when you create a version using "Save & Run All"
# You can also write temporary files to /kaggle/temp/, but they won't be saved outside of the current session
```

Let's load the dataset with the pandas framework and drop the unnecessary column. Then let's take a deeper look with what we are working here.

```
df = pd.read_csv('../input/sms-spam-collection-dataset/spam.csv', encoding = "ISO-8859-1")
df = df.drop(df.columns[range(2,5)], axis =1)
df.head()
```

Data Exporation and Analysis

First and foremost, let's check for any missing values in the dataset.

```
df_value_counts = df['v1'].value_counts()
df_value_counts
```

```
df_value_counts.isnull().sum()
```

We see that everything looks good and we don't have null or missing values. We are ready to move on.

Plotting the data

Firstly, let's get a sense of the amount of values within each category.

```
df_value_counts.plot.bar(color=['green', 'red'])
```

Secondly, I plot the 30 most common words within those categories.

```
from collections import Counter

ham_ct  = Counter(" ".join(df[df['v1']=='ham']["v2"]).split()).most_common(30)

spam_ct = Counter(" ".join(df[df['v1']=='spam']["v2"]).split()).most_common(30)

df_ham = pd.DataFrame.from_dict(ham_ct)

df_ham = df_ham.rename(columns={0: "word", 1 : "count"})

df_spam = pd.DataFrame.from_dict(spam_ct)

df_spam = df_spam.rename(columns={0: "word", 1 : "count"})


fig, axes = plt.subplots(1, 2, figsize=(12, 4))

df_ham.plot.bar(x='word', y='count', legend=False, ax=axes[0])

df_spam.plot.bar(x='word', y='count', color='red', legend=False, ax=axes[1])


axes[0].set_title('Non-Spam')

axes[1].set_title('Spam')


plt.show()
```

As we can see the most common words are for example prepositions, etc. These should not have any influence in the classification process.

Feature Engineering

This part, called *feature engineering*, is where we need to transform the raw dataset's data into usable features.

When dealing with text data from any source (email, SMS, etc) we need to *vectrize* this input into a representation called **Bag of words**.

This turns the input documents (corpus) into numerical representations, described by the amount of word occurrences.

As we saw earlier, these documents have various words that should have minimal influence, which are called *stop words*. We need to re-weight these

numerical representations into float values, which are more suitable into a classifier.

This technique is called **tf-idf transform** (term-frequency) and together with the occurences counter are combined in a single sklearn vectorizer.

```
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
transformer = TfidfVectorizer(stop_words = 'english')

X = transformer.fit_transform(df["v2"])

np.shape(X)
```

Data preparation

It is time to split the dataset into train and test sets. Before we do that, we should replace the words 'spam' and 'ham' into binary.

```
from sklearn.model_selection import train_test_split
```

```
# Extract features and prediction vectors
y = df['v1'].map({'ham':0, 'spam':1})

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 0)

X_train.shape, X_test.shape
```

Model creation and training

Let's initialize our two models, the MultinomialNB and the BernoulliNB. After initialization we fit the models to the data and record their scores for the prediction analysis.

```
from sklearn.naive_bayes import MultinomialNB, BernoulliNB


mnb = MultinomialNB()
mnb.fit(X_train, y_train)
score_train_mnb = mnb.score(X_train, y_train)
score_test_mnb  = mnb.score(X_test,  y_test)


bnb = BernoulliNB()
bnb.fit(X_train, y_train)
score_train_bnb = bnb.score(X_train, y_train)
score_test_bnb  = bnb.score(X_test,  y_test)
```

Prediction Analysis

The first thing we should do is to find what is the *null accuracy*. This value is the accuracy achieved by always selecting the most frequent class. To compute it we calculate the

occurences of the most frequent class, and divide by the total occurences of the test set. After that we go ahead and compute the mean accuracy for the train and test sets for both NB models.

```
print(f'{y_test.value_counts()}\n')


null_accuracy = (1434/(1434+238))
print('Null accuracy score: {0:0.4f}\n'.format(null_accuracy))


print("MultinomialNB:\n-> Train set: {}\n-> Test set:  {}".format(score_train_mnb, score_test_mnb))
print("BernoulliNB:\n-> Train set: {}\n-> Test set:  {}".format(score_train_bnb, score_test_bnb))
```

As we can see both models perform really well in classifing spam from non spam, with the Bernoulli model perfoming a little bit better. Also both model accuracies are

above the null accuracy value, which means they are good in classifing the specific task at hand.

Confusion Matrix

To further boost our confidence into these models we should use other metric tools for the underlying performance of the models. A well known tool is called the **confusion matrix**.

The confusion matrix is a 2×2 table that contains 4 outputs as a result of a binary classifier. There are four outcomes: TP, FP, TN, FN. In more detail:

* True Positives (TP)

* False Positives (FP)

* True Negatives (TN)

* False Negatives (FN)

from sklearn.metrics import confusion_matrix

y_pred_train_mnb = mnb.predict(X_train.toarray())

y_pred_train_bnb = bnb.predict(X_train.toarray())

cm_mnb_vals = confusion_matrix(y_train, y_pred_train_mnb)

cm_bnb_vals = confusion_matrix(y_train, y_pred_train_bnb)

cm_mnb = pd.DataFrame(data=cm_mnb_vals, columns=['Actual Positive (1)', 'Actual Negative (0)'],

                          index=['Predict Positive (1)', 'Predict Negative (0)'])

cm_bnb = pd.DataFrame(data=cm_bnb_vals, columns=['Actual Positive (1)', 'Actual Negative (0)'],

<div align="center">index=['Predict Positive (1)', 'Predict Negative (0)'])</div>

fig, axes = plt.subplots(1, 2, figsize=(12, 4))

axes[0].set_title('MultinomialNB')

axes[1].set_title('BernoulliNB')

sns.heatmap(cm_mnb, annot=True, fmt='d', ax=axes[0])

sns.heatmap(cm_bnb, annot=True, fmt='d', ax=axes[1])

plt.show()

Classification Measures

From the confusion matrix we can derive several classification measures:

1. Accuracy

This term tells us how many right predictions were made.

$\frac{TP+TN}{TP+FP+TN+FN}$

2. Error Rate

This term tells us how many wrong predictions were made.

$\frac{FP+FN}{TP+FP+TN+FN}$

3. Precision

Explains how many, of the predictions that were marked as positive, are actually truly positive.

$\frac{TP}{TP+FP}$

4. Sensitivity (Recall or True positive rate)

This term is a measure of how well the classifier can identify true positives.

$\frac{TP}{TP+FN}$

5. Specificity (True Negative rate)

This term is a measure of how well the classifier can identify true negatives.

$\frac{TN}{TN+FP}$

6. F1-Score (Harmonic mean of precision and recall)

This score can be interpreted as a weighted harmonic mean of the precision and recall. It maxes out when Recall = Precision.

$2*\frac{\text{Precision}*\text{Recall}}{\text{Precision}+\text{Recall}}$

```
# Classification measures for MultinomialNB
mnb_TP = cm_mnb_vals[0,0]
mnb_FP = cm_mnb_vals[0,1]
mnb_TN = cm_mnb_vals[1,1]
mnb_FN = cm_mnb_vals[1,0]


mnb_accuracy    = (mnb_TP + mnb_TN) / float(mnb_TP + mnb_TN + mnb_FP + mnb_FN)
mnb_error       = (mnb_FP + mnb_FN) / float(mnb_TP + mnb_TN + mnb_FP + mnb_FN)
mnb_precision   = mnb_TP / float(mnb_TP + mnb_FP)
mnb_recall      = mnb_TP / float(mnb_TP + mnb_FN)
mnb_specificity = mnb_TN / float(mnb_TN + mnb_FP)
```

```
mnb_f1        = 2 * (mnb_precision * mnb_recall) / (mnb_precision + mnb_recall)


print(f'MultinomialNB Accuracy:    {mnb_accuracy:.4f}')
print(f'MultinomialNB Error:       {mnb_error:.4f}')
print(f'MultinomialNB Precision:   {mnb_precision:.4f}')
print(f'MultinomialNB Recall:      {mnb_recall:.4f}')
print(f'MultinomialNB Specificity: {mnb_specificity:.4f}')
print(f'MultinomialNB F1-Score:    {mnb_f1:.4f}\n')


# Classification measures for BernoulliNB
bnb_TP = cm_bnb_vals[0,0]
bnb_FP = cm_bnb_vals[0,1]
bnb_TN = cm_bnb_vals[1,1]
bnb_FN = cm_bnb_vals[1,0]


bnb_accuracy    = (bnb_TP + bnb_TN) / float(bnb_TP + bnb_TN + bnb_FP + bnb_FN)
bnb_error       = (bnb_FP + bnb_FN) / float(bnb_TP + bnb_TN + bnb_FP + bnb_FN)
bnb_precision   = bnb_TP / float(bnb_TP + bnb_FP)
bnb_recall      = bnb_TP / float(bnb_TP + bnb_FN)
bnb_specificity = bnb_TN / float(bnb_TN + bnb_FP)
bnb_f1          = 2 * (bnb_precision * bnb_recall) / (bnb_precision + bnb_recall)


print(f'BernoulliNB Accuracy:    {bnb_accuracy:.4f}')
print(f'BernoulliNB Error:       {bnb_error:.4f}')
print(f'BernoulliNB Precision:   {bnb_precision:.4f}')
print(f'BernoulliNB Recall:      {bnb_recall:.4f}')
print(f'BernoulliNB Specificity: {bnb_specificity:.4f}')
print(f'BernoulliNB F1-Score:    {bnb_f1:.4f}')
```

As we can see from the metrics, the Bernoulli model behaves better in all aspects of the prediction process. It has better **accuracy** and lower **error** rate.

In addition the BernoulliNB has better **recall**, meaning it is better in identifying true positives, and as a result gives a better **F1-Score**.

ROC - AUC

Another very informative performance metric is called the *Receiver Operator Characteristic (ROC)*. It is a graphical plot that illustrates the diagnostic ability of a binary classifier system as its discrimination threshold is varied. [(Wikipedia)](https://en.wikipedia.org/wiki/Receiver_operating_characteristic)

The ROC curve is created by plotting the true positive rate (TPR - Sensitivity) against the false positive rate (FPR or (1 - Specificity)) at various threshold settings.

Firstly, let's find the probabilities for class 1 (Spam).

```
mnb_y_pred = mnb.predict_proba(X_test)[:, 1]
```

```
bnb_y_pred = bnb.predict_proba(X_test)[:, 1]
```

Calculate and plot the ROC curves for each model.

```
from sklearn.metrics import roc_curve
```

```
mnb_FPR, mnb_TPR, mnb_thresholds = roc_curve(y_test, mnb_y_pred)
```

```
bnb_FPR, bnb_TPR, bnb_thresholds = roc_curve(y_test, bnb_y_pred)
```

```
fig, axes = plt.subplots(1, 2, figsize=(12, 4))
```

```
axes[0].set_title('MultinomialNB')
```

```
axes[1].set_title('BernoulliNB')
```

```
axes[0].plot(mnb_FPR, mnb_TPR, linewidth=2)
```

```
axes[0].plot([0,1], [0,1], 'k--' )
```

```
axes[0].set_xlabel('False Positive Rate (1 - Specificity)')
```

```
axes[0].set_ylabel('True Positive Rate (Sensitivity)')
```

```
axes[1].plot(bnb_FPR, bnb_TPR, linewidth=2, color='y')
```

```
axes[1].plot([0,1], [0,1], 'k--' )

axes[1].set_xlabel('False Positive Rate (1 - Specificity)')

axes[1].set_ylabel('True Positive Rate (Sensitivity)')


plt.show()
```

As we can see the ROC curves are quite similar as expected. To accurately compare them we need to use the **Area Under Curve (AUC)** metric.

```
from sklearn.metrics import roc_auc_score


mnb_AUC = roc_auc_score(y_test, mnb_y_pred)

bnb_AUC = roc_auc_score(y_test, bnb_y_pred)


print(f'MultinomialNB AUC: {mnb_AUC:.4f}')

print(f'BernoulliNB AUC:   {bnb_AUC:.4f}')
```

Cross-Validation

Instead of using the pretty simple train-test dataset split, we should use another common practice to avoid *overfitting*. In order to better estimate the skill of our

classifiers to predict data on unseen data, we should use a technique called **k-fold Cross-Validation**. Cross validation is a resampling procedure, which splits the

data into k sets, and uses each set interchangably as the test set. It usually results to a less biased and less optimistic estimate of out models.

One popular value for k is the *10-fold Cross Validation*.

```
from sklearn.model_selection import cross_val_score


mnb_cv_accuracy = cross_val_score(mnb, X_train, y_train, cv = 10, scoring = 'accuracy')

mnb_cv_recall   = cross_val_score(mnb, X_train, y_train, cv = 10, scoring = 'recall')

mnb_cv_f1       = cross_val_score(mnb, X_train, y_train, cv = 10, scoring = 'f1')

mnb_cv_ROC      = cross_val_score(mnb, X_train, y_train, cv = 10, scoring = 'roc_auc')
```

```
print('MultinomialNB Cross-Validation:')

print(f'-> Accuracy: {mnb_cv_accuracy.mean():.4f}')

print(f'-> Recall:   {mnb_cv_recall.mean():.4f}')

print(f'-> F1-Score: {mnb_cv_f1.mean():.4f}')

print(f'-> ROC AUC:  {mnb_cv_ROC.mean():.4f}')


print()

bnb_cv_accuracy = cross_val_score(bnb, X_train, y_train, cv = 10, scoring = 'accuracy')

bnb_cv_recall   = cross_val_score(bnb, X_train, y_train, cv = 10, scoring = 'recall')

bnb_cv_f1       = cross_val_score(bnb, X_train, y_train, cv = 10, scoring = 'f1')

bnb_cv_ROC      = cross_val_score(bnb, X_train, y_train, cv = 10, scoring = 'roc_auc')

print('BernoulliNB Cross-Validation:')

print(f'-> Accuracy: {bnb_cv_accuracy.mean():.4f}')

print(f'-> Recall:   {bnb_cv_recall.mean():.4f}')

print(f'-> F1-Score: {bnb_cv_f1.mean():.4f}')

print(f'-> ROC AUC:  {bnb_cv_ROC.mean():.4f}')
```

Hyperparameter Tuning

There is still room for improvement. Both the Bernoulli and the Multinomial Naive Bayes models have a hyperparameter called **alpha** (α). This alpha is called the Laplace/Lidstone smoothing parameter.

Let's consider the classifier encounters a word which is not present in training set. The probability of existence of this word in any class is zero, which we need to avoid.

Here tuning the α parameter to high values makes the classifier biased towards the class with the most word occurences (underfitting).

In order to find a good value for α we can use techniques like *grid search* or *random search*. In this notebook I am going to use the more exhaustive sklearn's **GridSearchCV**, which is also using cross validation.

```
from sklearn.model_selection import GridSearchCV
```

```
param={'alpha': np.arange(0.01, 1, 0.01)}
```

```
mnb_grid_search = GridSearchCV(mnb, param, cv=10, scoring="accuracy",
return_train_score=True, n_jobs=-1)
```

```
mnb_grid_search.fit(X_train,y_train)
```

```
bnb_grid_search = GridSearchCV(bnb, param, cv=10, scoring="accuracy",
return_train_score=True, n_jobs=-1)
```

```
bnb_grid_search.fit(X_train,y_train)
```

```
mnb.set_params(alpha = mnb_grid_search.best_params_['alpha'])
```

```
bnb.set_params(alpha = bnb_grid_search.best_params_['alpha'])
```

```
print(f'MultinomialNB\n-> Best Alpha: {mnb_grid_search.best_params_["alpha"]}\n-> Mean
Accuracy Score: {cross_val_score(mnb, X_train, y_train, cv = 10,
scoring="accuracy").mean():.4f}')
```

```
print(f'BernoulliNB\n-> Best Alpha: {bnb_grid_search.best_params_["alpha"]}\n-> Mean
Accuracy Score: {cross_val_score(bnb, X_train, y_train, cv = 10,
scoring="accuracy").mean():.4f}')
```

Results

As we can see from all the performance metrics, both models behave very similarly to unseen data. We notice that the BernoulliNB model is better by a tiny margin in all scores.

We also saw that using cross-validation yields a better understanding of each model's predictive power, even though we didn't notice much of an accuracy increase compared to the simple train-test split.

We also performed a grid search in order to identify the best α hyperparameter for each model, boosting their predictive power.

I created this notebook to learn more about the naive bayes classification as well as how to carry out performance metrics on such models.

Any feedbacks and suggestions are mostly welcome!

Thank you.